

# Content Delivery Network Design

Bruno Zizi

February 8, 2024

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Preliminary Considerations and Assumptions . . . . .	2
<b>2</b>	<b>Objectives</b>	<b>2</b>
<b>3</b>	<b>Network Topology</b>	<b>3</b>
3.1	Master Nodes . . . . .	3
3.2	Edge Nodes . . . . .	3
3.3	DNS (Load Balancer) . . . . .	3
3.4	WAF (Web Application Firewall) . . . . .	3
3.5	Internal messaging system . . . . .	3
3.6	CDN Internal Network . . . . .	3
<b>4</b>	<b>Main Flows</b>	<b>5</b>
4.1	Bootstrap Procedure . . . . .	5
4.2	Image Spreading Procedure . . . . .	6
<b>5</b>	<b>Edge node implementation</b>	<b>8</b>
5.1	Communication with Master Node . . . . .	8
5.2	Bootstrap Procedure . . . . .	8
5.3	Caching Mechanism . . . . .	8
5.4	Logging and Tracing . . . . .	8
5.5	Health Check . . . . .	8
5.6	Additional Considerations . . . . .	9
<b>6</b>	<b>Master Node Implementation</b>	<b>11</b>
6.1	Handling HTTP Requests . . . . .	11
6.2	Image Spreading Procedure . . . . .	11
6.3	Bootstrap Procedure (Master Node) . . . . .	11
6.4	Maintaining Image Records . . . . .	11
6.5	Health Check Mechanism . . . . .	12
6.6	Logging and Tracing . . . . .	12
<b>7</b>	<b>Conclusion</b>	<b>13</b>
7.1	Consistency vs. Latency Trade-off . . . . .	13
7.2	Tailored Optimization . . . . .	13
7.3	Peer-to-Peer Communication . . . . .	13
7.4	Distributed Configuration . . . . .	13
7.5	Service Discovery . . . . .	13

# 1 Introduction

This document outlines the design of a Content Delivery Network (CDN) for the storage and retrieval of images through HTTP. Trying to adhere as much as possible to the specified requirements, the proposed solution aims to address the challenge of creating a distributed system within an on-premise infrastructure emphasizing speed, robustness, security, scalability, and maintainability.

## 1.1 Preliminary Considerations and Assumptions

- We value consistency more than latency in this implementation. We accept being slow in startup and during the Spread Procedure, but we can (almost) always sure that if an image is uploaded to the CDN, it will be available to all nodes.
- There is no one-size-fits-all solution. The implementation could be optimized in multiple ways.
- Granting a certain level of performance, always require a tailored solution. However multiple parameters are missing, such as the size of the CDN, expected traffic, security constraints (public vs. private), and so on.

## 2 Objectives

In this section of the document, we list a set of well-defined key objectives that would lead to an optimal solution for our CDN. These objectives guide the design and implementation of the CDN, ensuring that the solution is well-rounded, performs optimally, and meets the specified requirements.

1. **Speed:** Enable fast serving of content while also granting a fast and reliable spreading of newly uploaded images to all nodes in the network.
2. **Robustness:** Implement a resilient system that can withstand node failures and recover quickly.
3. **Security:** Mitigate DDoS attacks by consolidating Edge nodes behind a single IP using GBP protocol. Having a WAF in front of the Master nodes would also improve the overall security.
4. **Scalability:** Design the CDN to efficiently handle an increasing number of nodes and images.
5. **Maintainability:** Create a system that is easy to manage, ensuring the smooth operation of image storage and delivery.
6. **Availability:** Enable Edge nodes to efficiently serve images even if not yet synchronized locally.
7. **P2P Image Spreading:** Implement a P2P spreading mechanism for image distribution.
8. **Master Node Coordination:** Utilize Master nodes to coordinate image spreading, maintain image records, and perform health checks on nodes.
9. **Efficient Initial Synchronization (Edge node bootstrap):** Facilitate quick initial synchronization for Edge nodes joining the network.
10. **Edge Node Caching:** Leverage caching mechanisms whenever possible for improved performance.

## 3 Network Topology

The CDN is designed as a distributed system within an on-premise infrastructure. The topology involves multiple interconnected nodes each serving a specific purpose. Below is an overview of the key components and their relationships:

### 3.1 Master Nodes

Master Nodes play a central role in coordinating the CDN network. They maintain a comprehensive record of available images, orchestrate image spreading procedures, and perform health checks on connected Edge Nodes. The Master Nodes ensure the consistency and reliability of the CDN.

### 3.2 Edge Nodes

Edge Nodes act as the gateways for handling incoming HTTP requests and serving images. These nodes are distributed strategically to optimize content delivery and reduce latency. Each Edge Node communicates with the Master Node for synchronization and image spreading. From a very high level standpoint, they can be considered as some sort of Master Nodes distributed cache.

### 3.3 DNS (Load Balancer)

DNS nodes are responsible for directing incoming HTTP requests to the appropriate Edge Node. These components enhance the overall system's scalability, security and performance by efficiently distributing the incoming traffic.

### 3.4 WAF (Web Application Firewall)

The WAF adds an additional layer of security to the CDN infrastructure. Placed in front of the Master Nodes, it helps mitigate potential security threats, including Distributed Denial of Service (DDoS) attacks, thus ensuring availability.

### 3.5 Internal messaging system

Utilizing fast and asynchronous messaging systems such as RabbitMQ or Kafka could mitigate communication costs between nodes. This is achieved by minimizing the reliance on HTTP over TCP connections wherever possible.

### 3.6 CDN Internal Network

The CDN Internal Network comprises the interconnected Edge Nodes, Master Nodes, Load Balancer, and WAF. This internal network ensures secure communication and data transfer among CDN components while isolating them from the public internet.

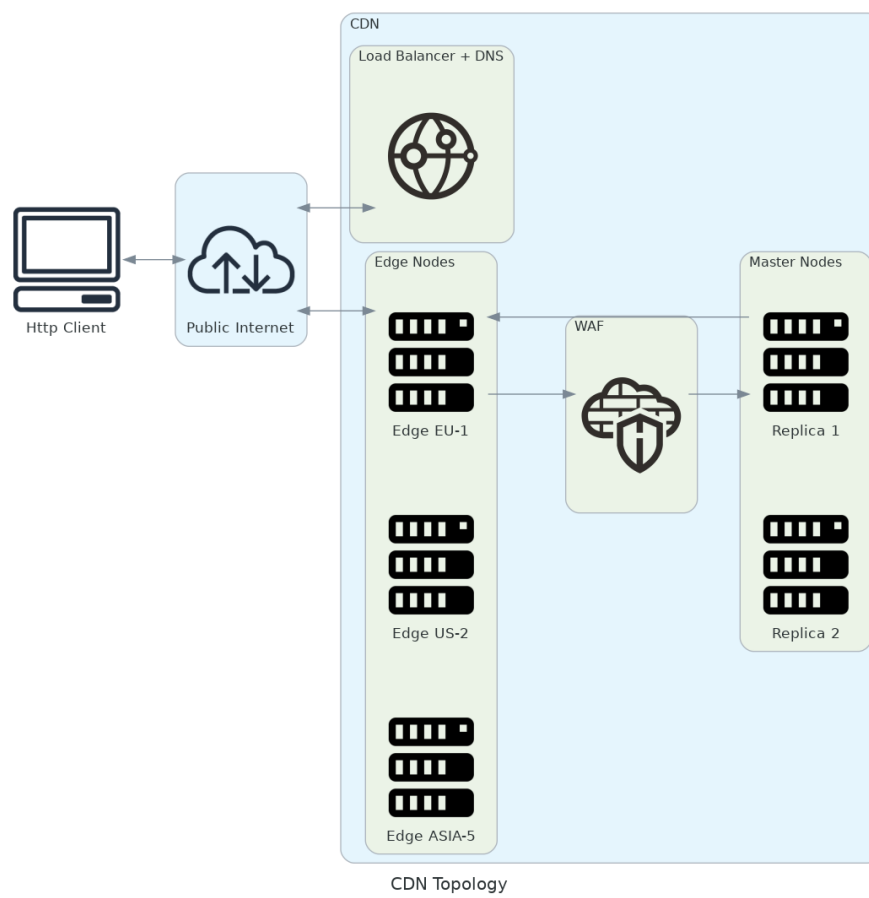


Figure 1: CDN Topology

## 4 Main Flows

### 4.1 Bootstrap Procedure

*Description:* The Bootstrap Procedure is initiated when an Edge Node joins the network for the first time or restarts after being down. Its purpose is to synchronize the new or recovering node with the existing network, ensuring it has the necessary information to serve images.

*Implementation:* The Edge Node communicates with the Master Node to obtain a list of available images and the next node(s) from which it should download images.

*Flow:*

1. When an Edge Node needs to bootstrap (either because it's joining the network for the first time or restarting after downtime), it sends a bootstrap request to the Master Node.
2. Upon receiving the bootstrap request, the Master Node responds with a map of available images and the corresponding IP addresses of nodes from which the images can be downloaded. If the bootstrap request contains a starting date, indicating a specific date from which the Edge Node needs images, only images updated after that date will be included.
3. The Edge Node, upon receiving the response from the Master Node, starts the process of downloading the images indicated in the map. To optimize the procedure, a 'batch download' operation could be implemented.
4. The Edge Node may also initiate a spread procedure to receive any new images uploaded to the network during its downtime.
5. Once the Bootstrap Procedure is complete, the Edge Node is synchronized with the network and can efficiently serve images to end-users.

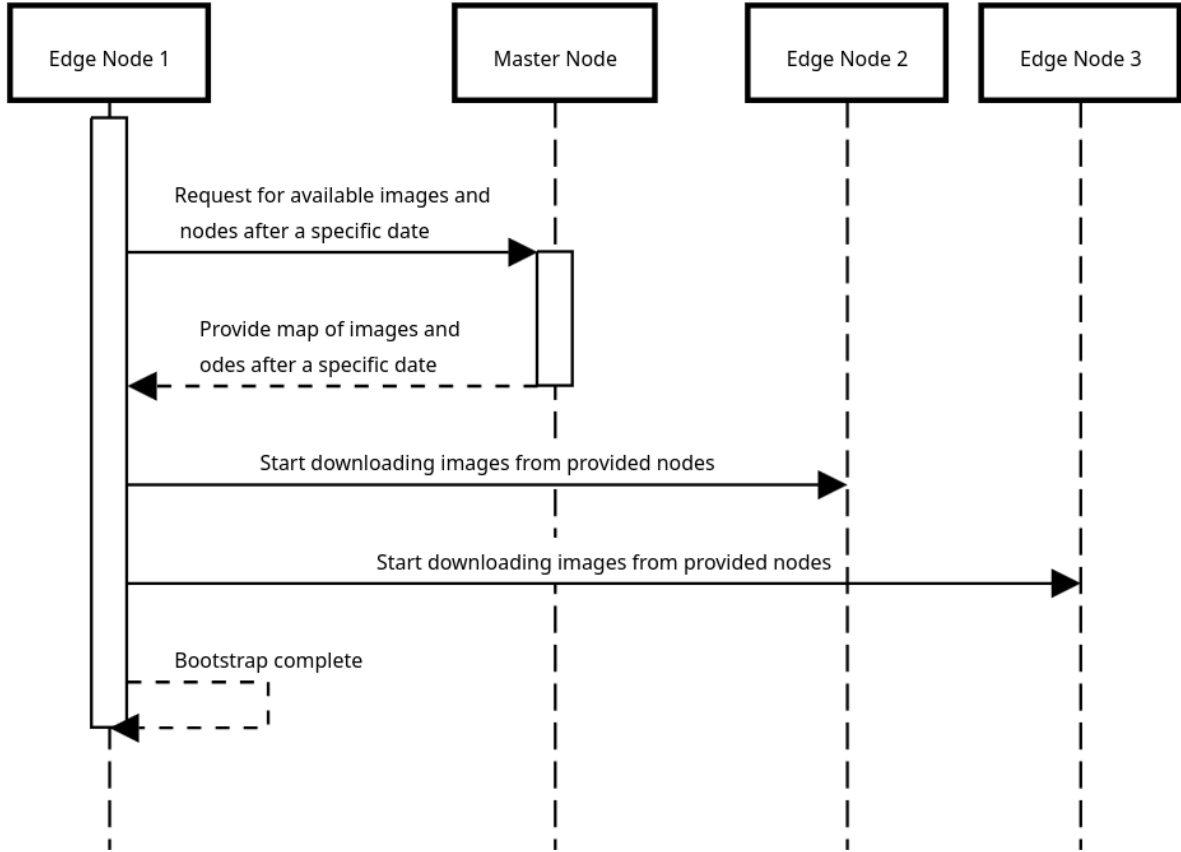


Figure 2: Bootstrap procedure

## 4.2 Image Spreading Procedure

*Description:* The Image Spread Procedure is initiated when a new image is uploaded to the network. Its purpose is to efficiently distribute the newly uploaded image to all connected Edge Nodes, ensuring quick availability for end-users.

*Implementation:* The Edge Node, upon receiving a new image upload request, notifies the Master Node about the new image. The Master Node, in turn, determines the next connected Edge Node(s) to which the image should be pushed for spreading.

*Flow:*

1. When a new image is uploaded, the Edge Node receiving the upload request adds the image to its list of available images and saves it locally.
2. The Edge Node notifies the Master Node about the new image upload, providing relevant details such as image metadata.
3. Upon receiving the notification, the Master Node updates its records to include the newly uploaded image in its list of available images.
4. The Master Node determines the next connected Edge Node(s) in the spreading sequence. This could be based on factors like load balancing, geographic proximity, or any defined distribution strategy.
5. The Master Node replies to the uploading Edge Node with the IP address(es) of the next node(s) in the spreading sequence.
6. The Edge node pushes the image to the next node.

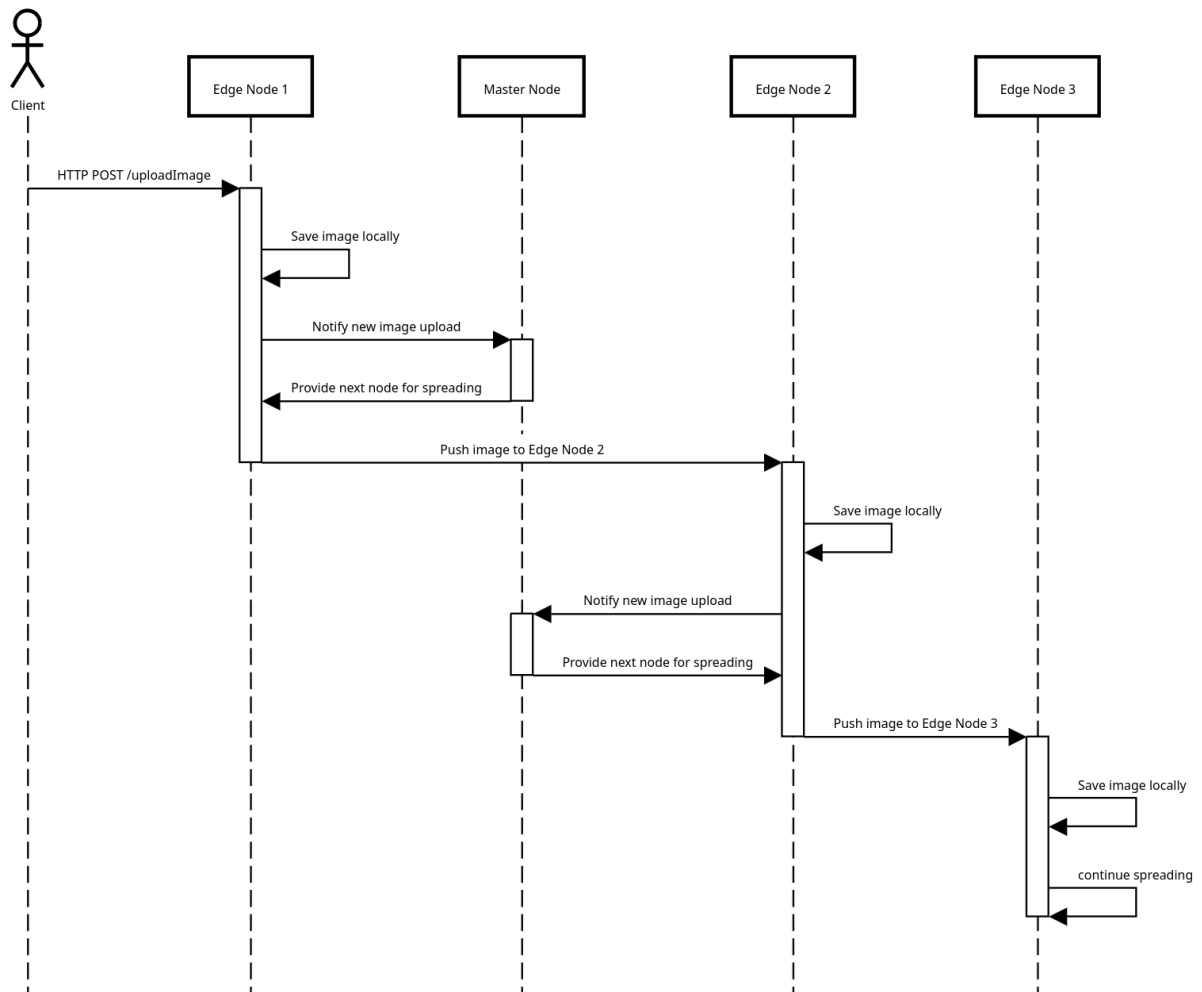


Figure 3: Image spreading procedure

## 5 Edge node implementation

### 5.1 Communication with Master Node

*Description:* Communicate with the Master Node for initial setup and the spread procedure.

*Implementation:* Use HTTP requests or a messaging system like Kafka or RabbitMQ for communication.

*Flow:*

1. Upon initialization, request the list of available images from the Master Node.
2. Utilize the spread procedure to receive all new images from the Master Node when they are uploaded.
3. If the node goes down and restarts, initiate a synchronization process with the Master Node.

### 5.2 Bootstrap Procedure

*Description:* Procedure for an Edge Node to join the network for the first time or after being down.

*Implementation:* Check with the Master Node for the list of available images and nodes for image downloading.

*Flow:*

1. When an Edge Node initiates a bootstrap request, the Master Node receives the request.
2. The Master Node responds with a map of available images and the corresponding IP addresses of nodes from which the images can be downloaded. If the bootstrap request includes a starting date, the Master Node tailors the map to contain only images updated after that specific date. This is useful in case of downtime recovery.
3. The Edge Node, upon receiving the response from the Master Node, starts downloading all the images indicated in the map.

### 5.3 Caching Mechanism

*Description:* Cache frequently accessed images locally for improved retrieval speed.

*Implementation:* Utilize a caching system like Redis or Memcached.

*Flow:*

1. Upon receiving an HTTP request, check the local cache for the requested image.
2. If the image is present and not expired, serve it directly.
3. If the image is not in the cache or has expired, fetch it from the Master Node and update the cache.

### 5.4 Logging and Tracing

*Description:* Record and trace relevant events for monitoring and debugging.

*Implementation:* Use a logging library and incorporate distributed tracing (e.g., OpenTelemetry).

*Flow:*

1. Log key events such as incoming requests, cache hits/misses, interactions with the Master Node and errors as well.

### 5.5 Health Check

*Description:* Regularly check the health of the Edge Node.

*Implementation:* Implement a health-check mechanism and leverage the Master Node's health-check capabilities.

*Flow:*

1. Expose an HTTP endpoint for health checks or use a dedicated Kafka topic.
2. Alternatively, an Edge Node could periodically ping the Master Node to report its health status.



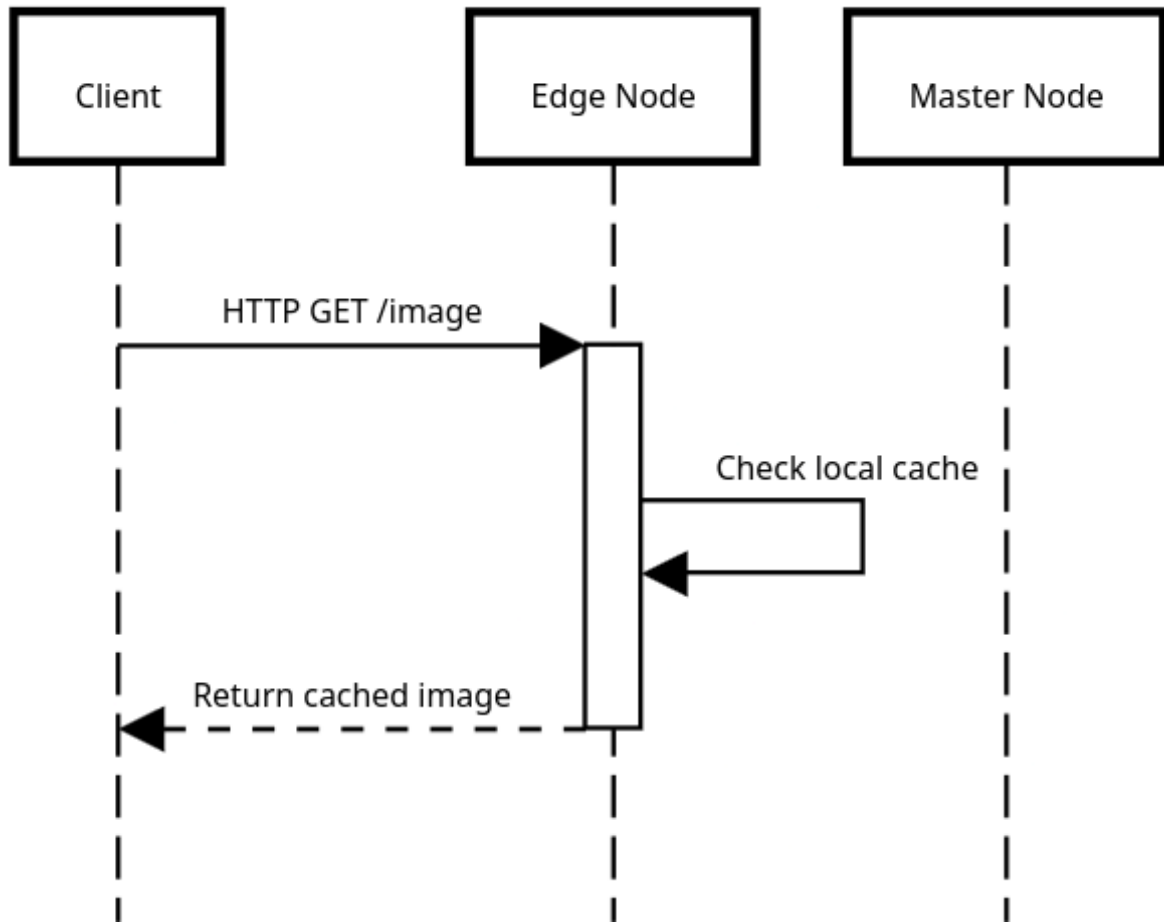


Figure 4: Edge node cache hit

## 5.6 Additional Considerations

### *Implementation:*

- Implement a configuration mechanism to adjust parameters such as cache expiration times and synchronization intervals.
- Employ exponential backoff strategies for retries during communication with the Master Node to handle potential network issues.




Figure 5: Edge cache miss

## 6 Master Node Implementation

### 6.1 Handling HTTP Requests

*Description:* Receive incoming HTTP requests for various operations (e.g., image spreading, health checks).

*Implementation:* Utilize a robust HTTP server library in C++ or Go.

*Flow:*

1. Listen for incoming HTTP requests.
2. Route requests based on operation types (spread, health check, etc.).
3. Handle each type of request accordingly.

### 6.2 Image Spreading Procedure

*Description:* Distribute newly uploaded images to all connected Edge Nodes.

*Implementation:* Utilize a push-based mechanism and reply with the IP(s) of the next node for spreading.

*Flow:*

1. Receive a new image upload request.
2. Add the image to the list of available images and save it locally.
3. Identify connected Edge Nodes that haven't already received the image.
4. Reply with the IP(s) of the next Edge Node(s) for spreading.

### 6.3 Bootstrap Procedure (Master Node)

*Description:* Procedure for an Edge Node to join the network for the first time or after being down.

*Implementation:* Check with the Master Node for the list of available images and nodes for image downloading.

*Flow:*

1. When an Edge Node needs to bootstrap (either because it's joining the network for the first time or restarting after downtime), it sends a bootstrap request to the Master Node.
2. Upon receiving the bootstrap request, the Master Node responds with a map of available images and the corresponding IP addresses of nodes from which the images can be downloaded. If the bootstrap request contains a starting date, indicating a specific date from which the Edge Node needs images, only images updated after that date will be included.
3. The Edge Node, upon receiving the response from the Master Node, starts the process of downloading the images indicated in the map. To optimize the procedure, a 'batch download' operation could be implemented.
4. The Edge Node may also initiate a spread procedure to receive any new images uploaded to the network during its downtime.
5. Once the Bootstrap Procedure is complete, the Edge Node is synchronized with the network and can efficiently serve images to end-users.

### 6.4 Maintaining Image Records

*Description:* Keep track of all available images in the CDN.

*Implementation:* An in-memory data store or a database can be used for the purpose.

*Flow:*

1. Maintain a data structure to store image metadata (e.g., filename, ID, timestamp).
2. Update the record upon each image upload.

## 6.5 Health Check Mechanism

*Description:* Receive (and handle) health checks data from connected Edge Nodes.

*Implementation:* Handle health-check requests from Edge Nodes. It can be implemented as a synchronous (HTTP) or an asynchronous (message queue) operation.

*Flow:*

1. Receive health checks data from connected Edge Nodes.
2. Update the health status for requesting nodes.

## 6.6 Logging and Tracing

*Description:* Record and trace relevant events for monitoring and debugging.

*Implementation:* Use a logging library and incorporate distributed tracing (e.g., OpenTelemetry).

*Flow:*

1. Log key events such as incoming requests, image spreading, health-check results as well as errors.

## 7 Conclusion

The designed CDN presents a comprehensive solution to the initial requirements, emphasizing consistency, speed, robustness, security, and scalability. However like any system, there are considerations regarding potential weak points and areas for improvement.

### 7.1 Consistency vs. Latency Trade-off

The chosen approach prioritizes consistency over latency during startup and image spreading procedures. This decision may result in slower initial synchronization and image spreading but guarantees that uploaded images are eventually available on all nodes. For scenarios where low-latency is critical, optimizing the trade-off may be explored based on specific use-case requirements.

### 7.2 Tailored Optimization

The design recognizes that there is no one-size-fits-all solution. Various parameters, such as CDN size, expected traffic, and security constraints, are missing or assumed. Tailoring the implementation based on these parameters could lead to further adjustments and optimization.

### 7.3 Peer-to-Peer Communication

While peer-to-peer communication enhances image spreading efficiency, challenges may arise in ensuring secure and reliable communication between Edge Nodes. Implementing robust mechanisms for peer discovery, authentication, and data integrity is essential to mitigate potential risks associated with peer-to-peer communication.

### 7.4 Distributed Configuration

One notable weak point in the current design is the lack of detailed considerations for a distributed configuration management system. As the CDN dynamically scales and grows, the need for an efficient configuration updates becomes crucial. Implementing a robust distributed configuration management system can enhance the flexibility and reduce the time to market significantly.

### 7.5 Service Discovery

Another area that requires attention is service discovery. The current design assumes a relatively static environment with predefined IP addresses for communication between nodes. In dynamic and large-scale deployments, a robust service discovery mechanism is essential to automatically locate and connect to nodes.